

Python

Python Concurrency

Part 3: Advanced Patterns and Hybrid Approaches

April 29, 2025



Source Code

1. Combining Approaches for Complex Systems


For real-world applications, you can combine different concurrency models to leverage their respective strengths:

```
1 import asyncio
2 import concurrent.futures
3 import time
4
5 def cpu_bound(number):
6     """CPU-bound task (runs in a process)"""
7     total = sum(i * i for i in range(number))
8     return total
9
10 def io_bound(number):
11     """I/O-bound task (runs in a thread)"""
12     time.sleep(1) # Simulate I/O
13     return number * 2
14
```



Alejandro Sánchez Yalí

Software Developer | AI & Blockchain Enthusiast


 www.asanchezyali.com

```
15 async def main():
16     # Create executor pools
17     process_pool =
        concurrent.futures.ProcessPoolExecutor(max_workers=4)
18     thread_pool =
        concurrent.futures.ThreadPoolExecutor(max_workers=10)
19
20     loop = asyncio.get_running_loop()
21
22     # CPU-bound tasks (run in process pool)
23     cpu_numbers = [5_000_000, 10_000_000, 15_000_000, 20_000_000]
24     cpu_tasks = [
25         loop.run_in_executor(process_pool, cpu_bound, number)
26         for number in cpu_numbers
27     ]
28
29     # I/O-bound tasks (run in thread pool)
30     io_numbers = list(range(1, 11))
31     io_tasks = [
32         loop.run_in_executor(thread_pool, io_bound, number)
```



Alejandro Sánchez Yalí

Software Developer | AI & Blockchain Enthusiast


 www.asanchezyali.com

```
33     for number in io_numbers
34 ]
35
36 # Async I/O-bound tasks (native asyncio)
37 async_tasks = [asyncio.sleep(1, result=f"async_{i}") for i in
range(5)]
38
39 # Gather all results
40 print("Running all tasks concurrently...")
41 start = time.time()
42
43 cpu_results = await asyncio.gather(*cpu_tasks)
44 io_results = await asyncio.gather(*io_tasks)
45 async_results = await asyncio.gather(*async_tasks)
46
47 end = time.time()
48
49 # Show results
50 print(f"\nTotal time: {end - start:.2f} seconds")
51 print(f"CPU results: {len(cpu_results)} tasks completed")
```



Alejandro Sánchez Yalí

Software Developer | AI & Blockchain Enthusiast

 www.asanchezyali.com

```
52     print(f"I/O results: {len(io_results)} tasks completed")
53     print(f"Async results: {len(async_results)} tasks completed")
54
55     # Clean up
56     process_pool.shutdown()
57     thread_pool.shutdown()
58
59 if __name__ == "__main__":
60     asyncio.run(main())
```

This hybrid approach demonstrates how to:

- **Use `ProcessPoolExecutor`** for CPU-bound tasks to bypass the GIL
- **Use `ThreadPoolExecutor`** for blocking I/O operations
- **Use native `asyncio`** for non-blocking I/O operations
- **Coordinate all approaches** through `asyncio`'s event loop

The example shows how `run_in_executor` allows integration of traditional concurrency approaches with `asyncio`, creating a unified system that handles different types of workloads optimally.



Alejandro Sánchez Yalí

Software Developer | AI & Blockchain Enthusiast

 www.asanchezyali.com

2. Advanced Concurrency Patterns

2.1. Fan-Out/Fan-In Pattern

The fan-out/fan-in pattern is ideal for data parallelism, where you split a large task into smaller subtasks, process them concurrently, and then combine the results:

```
1 import asyncio
2 import time
3 import random
4
5 async def fan_out_fan_in_example():
6     """
7     Demonstrates the fan-out/fan-in pattern:
8     1. Fan-out: Split a task into multiple subtasks
9     2. Process each subtask concurrently
10    3. Fan-in: Collect and combine results
11    """
12    async def process_chunk(chunk_id, data):
```



Alejandro Sánchez Yalí

Software Developer | AI & Blockchain Enthusiast


 www.asanchezyali.com

```
13     """Process a chunk of data"""
14     print(f"Processing chunk {chunk_id} with {len(data)} items")
15     await asyncio.sleep(random.uniform(0.5, 1.5)) # Simulate
processing
16
17     # Simulate results (sum of items with processing artifact)
18     result = sum(data) * random.uniform(0.9, 1.1)
19     print(f"Chunk {chunk_id} processed, result: {result:.2f}")
20     return result
21
22     # Create a large dataset
23     dataset = [i * 2 for i in range(1000)]
24
25     # 1. Fan-out: Split data into chunks
26     chunk_size = 100
27     chunks = [dataset[i:i+chunk_size] for i in range(0, len(dataset),
chunk_size)]
28     print(f"Split dataset into {len(chunks)} chunks of {chunk_size}
items each")
29
```



Alejandro Sánchez Yalí

Software Developer | AI & Blockchain Enthusiast


 www.asanchezyali.com

```
30     # 2. Process chunks concurrently
31     print("\nFanning out processing to multiple tasks...")
32     tasks = [
33         process_chunk(i, chunk)
34         for i, chunk in enumerate(chunks)
35     ]
36
37     # 3. Fan-in: Gather all results
38     print("\nFanning in results...")
39     start_time = time.time()
40     results = await asyncio.gather(*tasks)
41     end_time = time.time()
42
43     # Combine results (in this case, take the average)
44     final_result = sum(results) / len(results)
45
46     print(f"\nAll chunks processed in {end_time - start_time:.2f}
seconds")
47     print(f"Final result (average): {final_result:.2f}")
48
```



Alejandro Sánchez Yalí

Software Developer | AI & Blockchain Enthusiast

 www.asanchezyali.com


```
49 # Run the example
50 if __name__ == "__main__":
51     asyncio.run(fan_out_fan_in_example())
```

This pattern applies to many real-world scenarios like map-reduce operations, batch processing, and parallel data analysis.

2.2. Task Queue with Priority

A priority-based task queue allows processing important tasks first:

```
1 import asyncio
2 import random
3
4 class AsyncTaskQueue:
5     """A task queue with priority and worker pool for asyncio"""
6
7     def __init__(self, num_workers=3):
8         self.queue = asyncio.PriorityQueue()
```



Alejandro Sánchez Yalí


Software Developer | AI & Blockchain Enthusiast

 www.asanchezyali.com

```
9         self.num_workers = num_workers
10        self.workers = []
11        self.running = False
12
13    async def add_task(self, coro, priority=0):
14        """Add a task to the queue with priority (lower is higher)"""
15        await self.queue.put((priority, coro))
16
17    async def worker(self, worker_id):
18        """Worker that processes tasks from the queue"""
19        while self.running:
20            try:
21                # Get a task from the queue
22                priority, coro = await self.queue.get()
23
24                try:
25                    print(f"Worker {worker_id}: Processing task with
priority {priority}")
26
27                    # Execute the coroutine
28                    result = await coro
```

**Alejandro Sánchez Yalí**


Software Developer | AI & Blockchain Enthusiast

 www.asanchezyali.com

```
28             print(f"Worker {worker_id}: Task completed with
result: {result}")
29         except Exception as e:
30             print(f"Worker {worker_id}: Task failed with
error: {e}")
31         finally:
32             # Mark task as done
33             self.queue.task_done()
34         except asyncio.CancelledError:
35             break
36
37     async def start(self):
38         """Start the worker pool"""
39         self.running = True
40         self.workers = [
41             asyncio.create_task(self.worker(i))
42             for i in range(self.num_workers)
43         ]
44
45     async def stop(self):
```

**Alejandro Sánchez Yalí**


Software Developer | AI & Blockchain Enthusiast

 www.asanchezyali.com

```
46     """Stop the worker pool"""
47     self.running = False
48
49     # Wait for all tasks to complete
50     await self.queue.join()
51
52     # Cancel all workers
53     for worker in self.workers:
54         worker.cancel()
55
56     # Wait for all workers to complete cancellation
57     await asyncio.gather(*self.workers, return_exceptions=True)
58
59 # Example task implementations
60 async def data_processing_task(task_id, duration):
61     """Simulates a data processing task"""
62     await asyncio.sleep(duration) # Simulate work
63     return f>Data for task {task_id} processed"
64
65 async def demo_task_queue():
```

**Alejandro Sánchez Yalí**

Software Developer | AI & Blockchain Enthusiast


 www.asanchezyali.com

```
66     # Create a task queue
67     task_queue = AsyncTaskQueue(num_workers=3)
68
69     # Start the worker pool
70     await task_queue.start()
71
72     # Add tasks with different priorities
73     for i in range(10):
74         priority = random.randint(1, 3) # 1=high, 3=low priority
75         duration = random.uniform(0.5, 1.0)
76
77         # Create a task and add it to the queue
78         task = data_processing_task(i, duration)
79         await task_queue.add_task(task, priority)
80
81         print(f"Added Task {i} with priority {priority}")
82
83     # Wait for all tasks to complete and stop the worker pool
84     await task_queue.stop()
85     print("All tasks completed, worker pool stopped")
```



Alejandro Sánchez Yalí

Software Developer | AI & Blockchain Enthusiast

 www.asanchezyali.com

```
86
87 # Run the demo
88 if __name__ == "__main__":
89     asyncio.run(demo_task_queue())
```

This pattern is useful for building job queues, task schedulers, and work distribution systems.

3. Performance Considerations and Best Practices

3.1. Benchmarking Concurrent Code


When optimizing Python code with concurrency, it's essential to measure actual performance gains:

```
1 import time
2 import concurrent.futures
3
4 def benchmark(func, data, executor_class, max_workers=None):
5     """Benchmark a function using different execution methods"""
```



Alejandro Sánchez Yalí

Software Developer | AI & Blockchain Enthusiast


 www.asanchezyali.com

```
6     # Sequential execution (baseline)
7     start = time.time()
8     sequential_result = [func(item) for item in data]
9     sequential_time = time.time() - start
10    print(f"Sequential: {sequential_time:.4f}s")
11
12    # Concurrent execution
13    start = time.time()
14    with executor_class(max_workers=max_workers) as executor:
15        concurrent_result = list(executor.map(func, data))
16        concurrent_time = time.time() - start
17        print(f"Concurrent: {concurrent_time:.4f}s")
18        print(f"Speedup: {sequential_time/concurrent_time:.2f}x")
19
20    # Example CPU-bound task
21    def cpu_task(n):
22        """CPU-intensive calculation"""
23        return sum(i * i for i in range(n * 100000))
24
25    # Example I/O-bound task
```



Alejandro Sánchez Yalí

Software Developer | AI & Blockchain Enthusiast

 www.asanchezyali.com

```
26 def io_task(n):
27     """I/O-bound operation (simulated)"""
28     time.sleep(0.1) # Simulate I/O delay
29     return n * 2
30
31 # Demo for CPU-bound tasks
32 data = list(range(1, 9))
33 print("CPU-bound task with ProcessPoolExecutor:")
34 benchmark(cpu_task, data, concurrent.futures.ProcessPoolExecutor)
35
36 # Demo for I/O-bound tasks
37 print("\nI/O-bound task with ThreadPoolExecutor:")
38 benchmark(io_task, data, concurrent.futures.ThreadPoolExecutor)
```

3.2. Best Practices for Production Code

For production-grade concurrent Python applications, follow these guidelines:

- **Tool Selection:**

- I/O-bound → asyncio or threading (asyncio preferred for new code)
- CPU-bound → multiprocessing



Alejandro Sánchez Yalí

Software Developer | AI & Blockchain Enthusiast

 www.asanchezyali.com

- Mixed workloads → hybrid approach

- **Resource Management:**

- Reuse thread/process pools rather than creating new ones
- Use context managers or ensure proper cleanup in `finally` blocks
- Monitor memory usage, especially with multiprocessing

- **Error Handling:**

- Properly catch and handle exceptions in worker functions
- Use timeouts to prevent hanging operations
- Implement graceful shutdown mechanisms

- **Avoiding Common Pitfalls:**

- Thread Safety: Always protect shared resources with locks
- Deadlocks: Acquire locks in a consistent order
- Oversubscription: Don't create too many threads or processes


4. Conclusion (Part 3)

In this final part of our exploration of concurrency and parallelism in Python, we've covered advanced techniques that build on the fundamentals:



Alejandro Sánchez Yalí

Software Developer | AI & Blockchain Enthusiast

 www.asanchezyali.com

- We've learned how to combine different concurrency models (threading, multiprocessing, and asyncio) to create hybrid solutions that leverage the strengths of each approach.
- We've explored advanced concurrency patterns like the fan-out/fan-in pattern and priority-based task queues that solve real-world parallelization problems.
- We've examined practical benchmarking approaches to quantify performance improvements and make data-driven decisions.
- We've identified best practices and common pitfalls to create production-ready concurrent code.

By mastering these techniques, you can develop Python applications that efficiently utilize system resources, respond quickly to events, and process data in parallel. The key is selecting the right concurrency model for each specific task and combining them when needed for complex applications.

Remember that concurrency is not always the answer—sometimes a simpler sequential solution is more maintainable and even faster for small datasets. Always benchmark your code to ensure that your concurrent solution actually improves performance in your specific use case.



Alejandro Sánchez Yalí

Software Developer | AI & Blockchain Enthusiast



www.asanchezyali.com

5. References (Part 3)

- Python Documentation. *asyncio* - Asynchronous I/O. [Link](#)
- Python Documentation. *concurrent.futures* - Launching parallel tasks. [Link](#)
- Real Python. *Speed Up Your Python Program With Concurrency*. [Link](#)
- Brett Slatkin. *Effective Python*, Item 57-61: Concurrency and Parallelism. Addison-Wesley.
- Caleb Hattingh. *Using Asyncio in Python: Understanding Python's Asynchronous Programming Features*. O'Reilly Media.
- Translated, Edited and written in collaboration with AI.



Alejandro Sánchez Yalí

Software Developer | AI & Blockchain Enthusiast



www.asanchezyali.com

6. Explore My Other Posts

Enjoyed This Content?

Don't miss the previous part of this series:


Python Concurrency Part 2: Asyncio and HTTP Operations

Learn how to use Python's asyncio for efficient I/O operations, working with HTTP requests, and implementing producer-consumer patterns. Master the modern approach to concurrency.



Alejandro Sánchez Yalí

Software Developer | AI & Blockchain Enthusiast

 www.asanchezyali.com



Feedback

Found this **helpful?**

Save, comment and share

April 29, 2025